

# Database Systems - Homework 4

Yosef Goren & Yonatan Kahalani

July 6, 2023

## 1

### 1.1

```
1 for all tuples s in S:
2   for all tuples a in A:
3     if a.artistID == A.ID and A.genre == 'Jazz' and
4       s.releaseDate >= T0_DATE('2023/00/00', 'YYYY/MM/DD'):
5       output(a.ID, s.ID)
```

Let  $B_A$  be the number of blocks in  $A$ ,  $B_S$  the number of blocks in  $S$  and  $n_S$  the number of tuples in  $S$ .

As seen in the lectures the IO cost for this is  $O(B_A + n_S \cdot B_S) = O(10^4 + 10^7)$

### 1.2

- i. IO cost  $O(1)$ . By going through the tree and finding the lexicographically largest element with the constraint that `artistID=236363`, then getting the block which contains it - and extracting that element from that block.
- ii. Let  $n_S$  be the number of rows that refer to this artist, from the year 2023 onwards. Thus the IO complexity is  $O(n_S)$ .  
After going to the index and finding the largest leaf with `artistID = 236363`, we can generate the set of outputs (while outputting the `s.id` for each tuple) in decending order by iterating over the leaves which have a year after 2023.  
Note that the specific ordering is not required by the implementation is more efficient this way.

### 1.3

Iterate over each element from the set of artists, if the artist's genre is Jazz - output all of his songs which are from after 2023 - by searching through the tree in lexicographic order - in the same way as in ii. (Denote this alg with `alg12`), and finally we project the tuples on `A.ID` and `S.ID`.

```

1 for all tuples a in A:
2   if a.genre == 'Jazz':
3     //run alg from 1.2 ii running on a
4     T := alg12(a)
5     for t in T:
6       output(a.id, t.id)

```

## 2

### 2.1

- i. First we show the scheduling is not *view serializable* and thus it is not *conflict serializable* - thanks to the theorem seen in lectures which states that conflict serializability implies view serializability.

First we note that in any serial scheduling which is equivalent to the given scheduling,  $T_3$  has to come after  $T_2$  since  $T_3$  and  $T_2$  refer to the same variable  $y$ .

Thus the scheduling options remaining are:

- (a)  $T_1 \rightarrow T_2 \rightarrow T_3$ .  
This scheduling is not equivalent since in the original schedule -  $T_1$  reads  $y$  written by  $T_2$  and in this one it reads the initial value of  $y$ .
- (b)  $T_2 \rightarrow T_3 \rightarrow T_1$   
Not equivalent since in the original the final value of  $x$  is written by  $T_2$  while in this schedule it is written by  $T_1$ .
- (c)  $T_2 \rightarrow T_1 \rightarrow T_3$   
Not equivalent due to exact the same reason as  $T_2 \rightarrow T_3 \rightarrow T_1$ .

- ii. Denote the following serialization:

$$S' = R_2(z), R_2(y), W_2(x), W_3(z), W_4(z), R_4(y), R_1(z)$$

The operations are equivalent:  $S =_C S'$ , since each pair of operations in conflict in  $S$  appear in the same relative order in  $S'$ .

For example the conflicting operations  $R_2(z), W_3(z)$ , are in the ordering  $R_2(z) \rightarrow \dots \rightarrow W_3(z)$  both in the original and new schedule.

Thus  $S$  is *conflict serializable* by def.

Thus from the theorem seen in lecture -  $S$  is also *view serializable*.

- iii. Denote the provided schedule as  $S$ .  
Note that  $T_2$  has to appear last in any serial scheduling which is equivalent to the one provided, since all transactions write to  $y$ , while  $W_2(y)$  happens after both  $W_1(y)$  (meaning  $T_2$  must be after  $T_1$ ) and  $W_2(y)$  is after  $W_3(y)$  (meaning  $T_2$  must be after  $T_3$ ).  
Thus we are left with two options:

- (a)  $T_1 \rightarrow T_3 \rightarrow T_2$
- (b)  $T_3 \rightarrow T_1 \rightarrow T_2$

$b$  is not *view equivalent* to  $S$  since  $S$ ,  $T_1$  reads  $y$ 's initial value, while in  $b$  -  $T_1$  reads the value written to  $y$  in  $T_3$ .

$a$  is *view equivalent* to  $S$  by def. and thus  $S$  is *view serializable*.

So far we have found that  $a$  is the only possible equivalent serialization of  $S$ . Since in  $S$ :  $W_3(y) \rightarrow \dots \rightarrow W_1(y)$ , and in  $a$ :  $W_1(y) \rightarrow \dots \rightarrow W_3(y)$ , thus  $a$  is not *conflict equivalent* to  $S$ .

This means there is no *conflict equivalent* serialization for  $S$ .

Hence  $S$  is not *conflict serializable*.

## 2.2

- i.  $R_{L_1}(x), R_1(x), R_{L_2}(y), R_2(y), R_1(x), R_{U_1}(x), R_{L_2}(x), R_2(x), R_2(y), R_{U_2}(x), R_{U_2}(y)$
- ii.  $W_2(x), W_1(x), W_2(x), R_2(y)$

## 3

For the purpose of providing an example for each solution we have provided in this question we have created an example database and added a few kittens to it, and ran the solution queries on each of them.

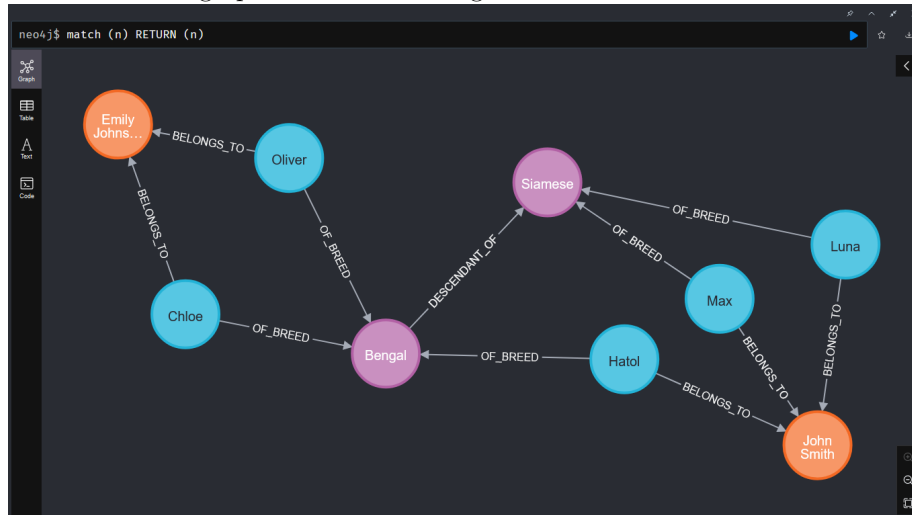
The following code creates the content of our database:

```

1 CREATE (b1:Breed {name: 'Siamese', origin: 'Thailand', lifespan:
   '12-15 years'})
2 CREATE (b2:Breed {name: 'Bengal', origin: 'United States', lifespan
   : '12-16 years'})
3 CREATE (b2)-[:DESCENDANT_OF]->(b1)
4
5 CREATE (o1:Owner {name: 'John Smith', address: '123 Main St'})
6 CREATE (o2:Owner {name: 'Emily Johnson', address: '456 Elm St'})
7
8 CREATE (c1:Cat {name: 'Max', age: 3, gender: 'Male'})
9 CREATE (c2:Cat {name: 'Luna', age: 2, gender: 'Female'})
10 CREATE (c3:Cat {name: 'Oliver', age: 4, gender: 'Male'})
11 CREATE (c4:Cat {name: 'Chloe', age: 1, gender: 'Female'})
12 CREATE (c5:Cat {name: 'Hatol', age:7, gender: 'Male'})
13
14 CREATE (c1)-[:BELONGS_TO {relationship_length: 2}]->(o1)
15 CREATE (c2)-[:BELONGS_TO {relationship_length: 1}]->(o1)
16 CREATE (c3)-[:BELONGS_TO {relationship_length: 1}]->(o2)
17 CREATE (c4)-[:BELONGS_TO {relationship_length: 1}]->(o2)
18 CREATE (c5)-[:BELONGS_TO {relationship_length: 3}]->(o1)
19
20 CREATE (c1)-[:OF_BREED]->(b1)
21 CREATE (c2)-[:OF_BREED]->(b1)
22 CREATE (c3)-[:OF_BREED]->(b2)
23 CREATE (c4)-[:OF_BREED]->(b2)
24 CREATE (c5)-[:OF_BREED]->(b2)

```

Here is how the graph looks after adding the kittens:



### 3.1

The query can be solved by:

```
1 MATCH (b:Breed)
2 OPTIONAL MATCH (b)-[:OF_BREED]-(c:Cat)
3 RETURN b.name, count(c)
```

Here is what it yields on our example:

b.name	count(c)
"Siamese"	2
"Bengal"	3

Started streaming 2 records after 2 ms and completed after 3 ms.

Running this on our example yields:

<b>b.name</b>	<b>count(c)</b>
"Siamese"	2
"Bengal"	3

### 3.2

This question is solved by the query:

```

1 MATCH (o:Owner) <-[:BELONGS_TO]-(c:Cat)-[:OF_BREED]->(b:Breed)
2 WITH o, collect(DISTINCT b) AS breeds
3 WHERE size(breeds) = 1
4 RETURN o.name

```

Running this on our example yields:

<b>b.name</b>	<b>count(DISTINCT c)</b>
"Siamese"	5
"Bengal"	3

### 3.3

This question is solved by the query:

```

1 MATCH (b:Breed)
2 OPTIONAL MATCH (b) <-[:DESCENDANT_OF*0..]-(c:Breed)-[:OF_BREED]-(c)
3 RETURN b.name, count(DISTINCT c)

```

Running this on our example yields:

<b>o.name</b>
"Emily Johnson"

### 3.4

This question is solved by the query:

```

1 MATCH (b:Breed)
2 OPTIONAL MATCH (b) <-[:OF_BREED]-(c:Cat)
3 RETURN b.name, count(c)

```

Running this on our example yields:

<b>o.name</b>
"John Smith"